



EFFECTIVE WEB SERVICE LOAD BALANCING THROUGH STATISTICAL MONITORING

By GEORGE PORTER *and* RANDY H. KATZ

W

Identifying the correlated effects between components to improve response to overload.

Web services are increasingly used for deploying Web-based application portals such as eBay and Amazon.com. They are commonly built using middleware, that is, composable building blocks such as HTTP containers, Java-based application beans, and persistent state management. These components are distributed across tiers of servers—Web, application, and database. As Web services offer newer and more sophisticated functionality, the underlying realization of those services in the middleware becomes more complicated. Today's Web services can consist of dozens or hundreds of request types and middleware components.

ILLUSTRATION BY JEAN-FRANÇOIS PODEVIN